Identifiability

Joachim Vandekerckhove
Spring 2025

Recap: What is Non-Identifiability?

A model is non-identifiable if different parameter combinations yield the exact same predictions (likelihood).

Recap: What is Non-Identifiability?

A model is non-identifiable if different parameter combinations yield the exact same predictions (likelihood).

The data cannot distinguish between these combinations.

Recap: What is Non-Identifiability?

A model is non-identifiable if different parameter combinations yield the exact same predictions (likelihood).

The data cannot distinguish between these combinations.

Structural Non-Identifiability: Inherent in model equations.

Recap: What is Non-Identifiability?

A model is non-identifiable if different parameter combinations yield the exact same predictions (likelihood).

The data cannot distinguish between these combinations.

Structural Non-Identifiability: Inherent in model equations.

Consequences:

- Poor MCMC convergence ($\hat{R} \gg 1$, low ESS).
- Unreliable posterior distributions (drifting, odd shapes).
- Strong correlations between confounded parameters.
- Misleading inferences.

1

Example: Hierarchical SDT Model

Predicting Hits (H_{ip}) and False Alarms (F_{ip}) for person p under condition i.

Example: Hierarchical SDT Model

Predicting Hits (H_{ip}) and False Alarms (F_{ip}) for person p under condition i.
Likelihood: Binomial counts based on SDT probabilities.

$$(H_{ip}, F_{ip}) \sim SDT\left(d'_{ip}, c_p, n_s, n_n\right)$$

Example: Hierarchical SDT Model

Predicting Hits (H_{ip}) and False Alarms (F_{ip}) for person p under condition i.
Likelihood: Binomial counts based on SDT probabilities.

$$(H_{ip}, F_{ip}) \sim SDT\left(d'_{ip}, c_p, n_s, n_n\right)$$

Hierarchical Parameter Models (Non-Identified Version):

$$
\begin{aligned}
d'_{ip} &= \alpha_p + \gamma_i & \\
\gamma_i &= \zeta_0 + \zeta_1 \cdot \text{predictor}_i & \text{(Condition effect on } d') \\
\alpha_p &\sim \text{Normal}(\mu_\alpha, \sigma_\alpha^2) & \text{(Person sensitivity intercept)} \\
c_p &\sim \text{Normal}(\mu_c, \sigma_c^2) & \text{(Person criterion)}
\end{aligned}
$$

Spotting the Non-Identifiability

The core issue lies in the definition of sensitivity d'_{ip}:

$$\begin{aligned}
d'_{ip} &= \alpha_p + \gamma_i \\
&= \alpha_p + (\zeta_0 + \zeta_1 \cdot \mathsf{predictor}_i) \\
&= (\alpha_p + \zeta_0) + \zeta_1 \cdot \mathsf{predictor}_i
\end{aligned}$$

Spotting the Non-Identifiability

The core issue lies in the definition of sensitivity d'_{ip}:

$$\begin{aligned} d'_{ip} &= \alpha_p + \gamma_i \\ &= \alpha_p + (\zeta_0 + \zeta_1 \cdot \mathsf{predictor}_i) \\ &= (\alpha_p + \zeta_0) + \zeta_1 \cdot \mathsf{predictor}_i \end{aligned}$$

The likelihood only depends on the sum $(\alpha_p + \zeta_0)$.

Spotting the Non-Identifiability

The core issue lies in the definition of sensitivity d'_{ip}:

$$d'_{ip} = \alpha_p + \gamma_i$$
$$= \alpha_p + (\zeta_0 + \zeta_1 \cdot \mathsf{predictor}_i)$$
$$= (\alpha_p + \zeta_0) + \zeta_1 \cdot \mathsf{predictor}_i$$

The likelihood only depends on the sum $(\alpha_p + \zeta_0)$.

Let $\tilde{\alpha}_p = \alpha_p - C$ and $\tilde{\zeta}_0 = \zeta_0 + C$.

Then $\tilde{\alpha}_p + \tilde{\zeta}_0 = \alpha_p + \zeta_0$. The predicted d'_{ip} is unchanged.

Spotting the Non-Identifiability

The core issue lies in the definition of sensitivity d'_{ip}:

$$\begin{aligned} d'_{ip} &= \alpha_p + \gamma_i \\ &= \alpha_p + (\zeta_0 + \zeta_1 \cdot \text{predictor}_i) \\ &= (\alpha_p + \zeta_0) + \zeta_1 \cdot \text{predictor}_i \end{aligned}$$

The likelihood only depends on the sum $(\alpha_p + \zeta_0)$.

Let $\tilde{\alpha}_p = \alpha_p - C$ and $\tilde{\zeta}_0 = \zeta_0 + C$.

Then $\tilde{\alpha}_p + \tilde{\zeta}_0 = \alpha_p + \zeta_0$. The predicted d'_{ip} is unchanged.

As a result, there cannot exist a unique set of optimal parameters.

Spotting the Non-Identifiability

The core issue lies in the definition of sensitivity d'_{ip}:

$$d'_{ip} = \alpha_p + \gamma_i$$
$$= \alpha_p + (\zeta_0 + \zeta_1 \cdot \mathsf{predictor}_i)$$
$$= (\alpha_p + \zeta_0) + \zeta_1 \cdot \mathsf{predictor}_i$$

The likelihood only depends on the sum $(\alpha_p + \zeta_0)$.

Let $\tilde{\alpha}_p = \alpha_p - C$ and $\tilde{\zeta}_0 = \zeta_0 + C$.

Then $\tilde{\alpha}_p + \tilde{\zeta}_0 = \alpha_p + \zeta_0$. The predicted d'_{ip} is unchanged.

As a result, there cannot exist a unique set of optimal parameters.

Nothing can distinguish trade-offs between μ_α and ζ_0.

Spotting the Non-Identifiability

The core issue lies in the definition of sensitivity d'_{ip}:

$$d'_{ip} = \alpha_p + \gamma_i$$
$$= \alpha_p + (\zeta_0 + \zeta_1 \cdot \mathsf{predictor}_i)$$
$$= (\alpha_p + \zeta_0) + \zeta_1 \cdot \mathsf{predictor}_i$$

The likelihood only depends on the sum $(\alpha_p + \zeta_0)$.

Let $\tilde{\alpha}_p = \alpha_p - C$ and $\tilde{\zeta}_0 = \zeta_0 + C$.

Then $\tilde{\alpha}_p + \tilde{\zeta}_0 = \alpha_p + \zeta_0$. The predicted d'_{ip} is unchanged.

As a result, there cannot exist a unique set of optimal parameters.

Nothing can distinguish trade-offs between μ_α and ζ_0.

This is structural non-identifiability.

Common MCMC Symptoms

Fitting the non-identified model is expected to show:

Common MCMC Symptoms

Fitting the non-identified model is expected to show:

– **Poor Convergence** for confounded parameters (μ_α, ζ_0): High $\hat{R}$, Low ESS.

Common MCMC Symptoms

Fitting the non-identified model is expected to show:

– **Poor Convergence** for confounded parameters (μ_α, ζ_0): High $\hat{R}$, Low ESS.

– **Wandering Trace Plots**: Chains for μ_α, ζ_0 may not mix or stabilize.

Common MCMC Symptoms

Fitting the non-identified model is expected to show:

– **Poor Convergence** for confounded parameters (μ_α, ζ_0): High $\hat{R}$, Low ESS.

– **Wandering Trace Plots**: Chains for μ_α, ζ_0 may not mix or stabilize.

– **Strong Posterior Correlation**: A clear negative correlation between μ_α and ζ_0 samples.

Common MCMC Symptoms

Fitting the non-identified model is expected to show:

- **Poor Convergence** for confounded parameters (μ_α, ζ_0): High $\hat{R}$, Low ESS.

- **Wandering Trace Plots**: Chains for μ_α, ζ_0 may not mix or stabilize.

- **Strong Posterior Correlation**: A clear negative correlation between μ_α and ζ_0 samples.

- **Sampling Difficulties**: Potential warnings (divergences, etc.).

Any or all of these can occur.

PyMC Code (Non-Identified SDT)

Key parameter definitions causing the issue:

```
1 # Uninformative priors for Sensitivity (d') parameters
2 mu_alpha = pm.Normal("mu_alpha", mu=0.0, sigma=1.0e9) # << Part 1
3 sigma_alpha = pm.HalfNormal("sigma_alpha", sigma=1.0e9)
4 zeta0 = pm.Normal("zeta0", mu=0.0, sigma=1.0e9)          # << Part 2
5 zeta1 = pm.Normal("zeta1", mu=0.0, sigma=1.0e9)
6
7 # Person sensitivity intercepts (drawn from mu_alpha)
8 alpha_p = pm.Normal("alpha_p", ...)
9
10 # Calculate d' using alpha_p and gamma_i
11 d_ip = pm.Deterministic("d_ip", alpha_p[...] + gamma_i) # << Part 3
12
13 # ... SDT calculations ...
```

Possible Diagnosis 1: Trace Plots

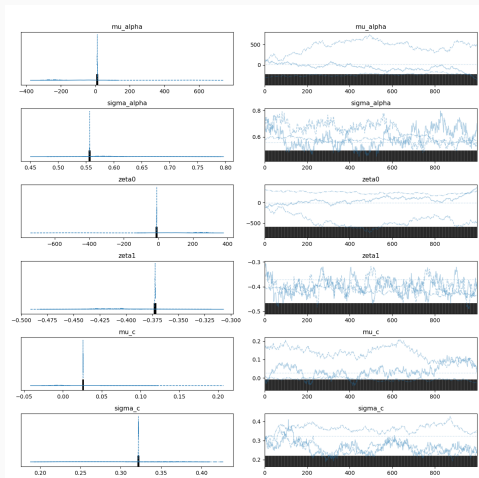Use 'arviz.plot_trace()' focusing on 'mu_alpha' and 'zeta0'.

Possible Diagnosis 1: Trace Plots

Use 'arviz.plot_trace()' focusing on 'mu_alpha' and 'zeta0'.

Look For:
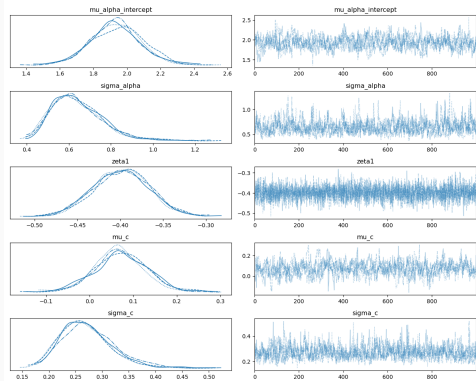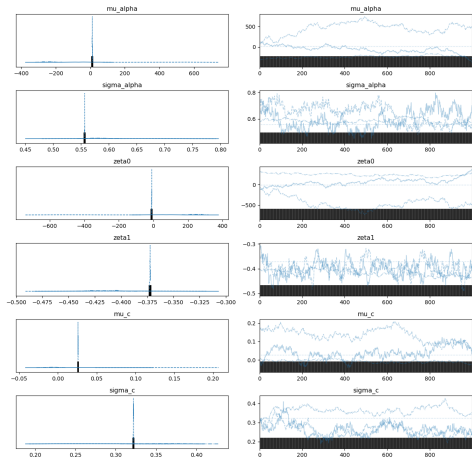
- Poor mixing for 'mu_alpha', 'zeta0'.

- Chains exploring different levels.

- Compare to well-behaved traces (e.g., 'mu_c').

- Corroborate with high $\hat{R}$ in summary.

For comparison, here are trace plots of the fixed (identified) model:

MCMC Output Table

```
    --- Non-Identified SDT Model Summary ---
                     mean       sd   hdi_3%  hdi_97%  mcse_mean  mcse_sd  ess_bulk  ess_tail  r_hat
mu_alpha           40.779  268.443 -304.985  570.827    128.271   70.199       5.0      17.0   2.70
zeta0             -38.809  268.503 -569.032  307.068    128.298   70.209       5.0      17.0   2.70
zeta1              -0.397    0.028   -0.446   -0.350      0.008    0.001      15.0     180.0   1.22
mu_c                0.048    0.058   -0.020    0.173      0.027    0.014       5.0      26.0   2.35
Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 379 seconds.
There were 1000 divergences after tuning. Increase 'target_accept' or reparameterize.


    --- Identified SDT Model Summary ---
                       mean     sd   hdi_3%  hdi_97%  mcse_mean  mcse_sd  ess_bulk  ess_tail  r_hat
mu_alpha_intercept    1.936  0.149    1.659    2.221      0.008    0.004     380.0     598.0   1.01
zeta1                -0.398  0.033   -0.463   -0.341      0.001    0.000    3558.0    3080.0   1.00
mu_c                  0.072  0.062   -0.046    0.186      0.003    0.002     426.0     735.0   1.01
Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 12 seconds.
```
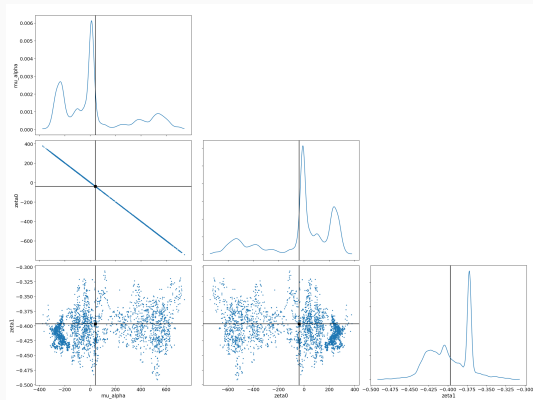
Possible Diagnosis 2: Pair Plots

Use 'arviz.plot_pair()' for sensitivity parameters 'mu_alpha', 'zeta0'.

Possible Diagnosis 2: Pair Plots

Use 'arviz.plot_pair()' for sensitivity parameters 'mu_alpha', 'zeta0'.

Look For:

- Strong linear correlation "ridge."

- Specifically, a negative correlation between 'mu_alpha' and 'zeta0'.

- Poorly conditioned marginal densities.

Fixing the Non-Identifiability

Common strategies:

Fixing the Non-Identifiability

Common strategies:

a) Reparameterization (Preferred):

- Remove redundant parameters (e.g., remove ζ_0).
- Use sum-to-zero constraints (more complex).
- Use contrasts relative to a baseline condition.

Fixing the Non-Identifiability

Common strategies:

a) Reparameterization (Preferred):

– Remove redundant parameters (e.g., remove ζ_0).

– Use sum-to-zero constraints (more complex).

– Use contrasts relative to a baseline condition.

b) Informative Priors: Can sometimes help practically, but remember that your conclusions will then strongly depend on your assumptions.

Fixing the Non-Identifiability

Common strategies:

a) Reparameterization (Preferred):

- Remove redundant parameters (e.g., remove ζ_0).
- Use sum-to-zero constraints (more complex).
- Use contrasts relative to a baseline condition.

b) Informative Priors: Can sometimes help practically, but remember that your conclusions will then strongly depend on your assumptions.

Let's see what informative priors can do...

Using Informative Priors

Setting $\mu_\alpha \sim N(0, 1.5)$ gives converging chains:

```
    --- Semi-Identified SDT Model Summary ---
        mean     sd  hdi_3%  hdi_97%  mcse_mean  mcse_sd  ess_bulk  ess_tail  r_hat
mu_alpha   1.301  0.827  -0.310    2.807      0.018    0.013    2023.0    2365.0   1.00
zeta0      0.629  0.825  -0.924    2.183      0.018    0.013    2002.0    2339.0   1.00
zeta1     -0.398  0.033  -0.459   -0.338      0.001    0.001    3915.0    2744.0   1.00
mu_c       0.067  0.064  -0.051    0.190      0.003    0.002     480.0     720.0   1.01
```
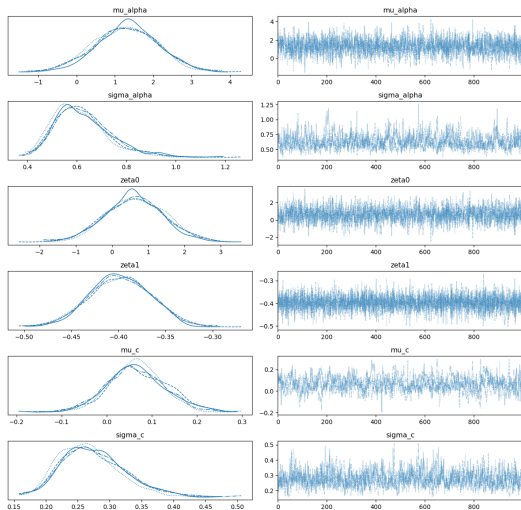
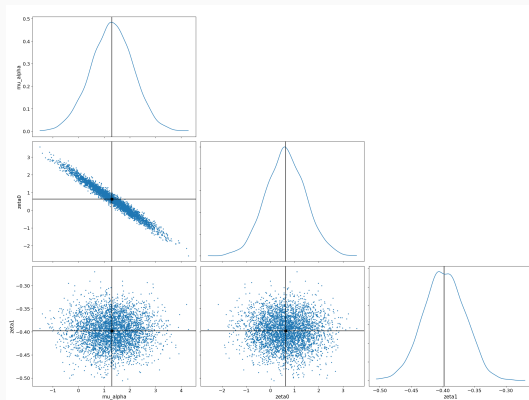## Using Informative Priors

Even reasonable trace plots:
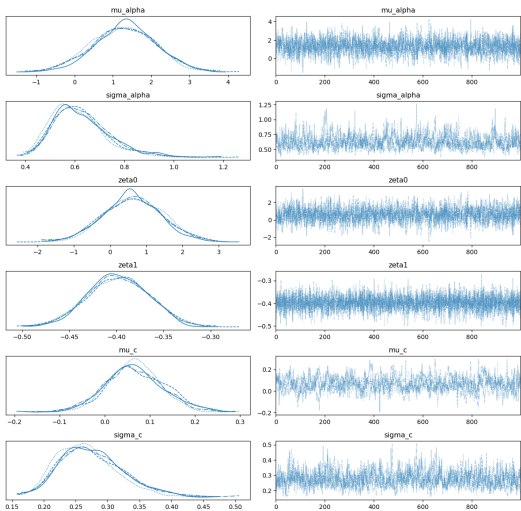
Using Informative Priors

Even reasonable trace plots:

But the trade-off pops out of the pair plot:

Applying the Real Fix: Remove ζ_0

Interpretation changes: μ_α becomes 'mu_alpha_int', average d' when predictor=0.

Applying the Real Fix: Remove ζ_0

Interpretation changes: μ_α becomes 'mu_alpha_int', average d' when predictor=0.

```
1  # Priors for Sensitivity (d') parameters
2  mu_alpha_int = pm.Normal("mu_alpha_int", ...) # Was mu_alpha
3  sigma_alpha = pm.HalfNormal("sigma_alpha", ...)
4  # zeta0 = pm.Normal("zeta0", ...) # <<< REMOVED
5  zeta1 = pm.Normal("zeta1", ...)
6
7  # Person sensitivity intercepts (drawn from mu_alpha_intercept)
8  alpha_p = pm.Normal("alpha_p", ...)
9
10 # Condition Effects on Sensitivity (NO intercept)
11 gamma_i = pm.Deterministic("gamma_i", zeta1 * cond_pred_data)
12
13 # SDT likelihood calculations are structurally identical
14 # ...
```

Results After Fixing
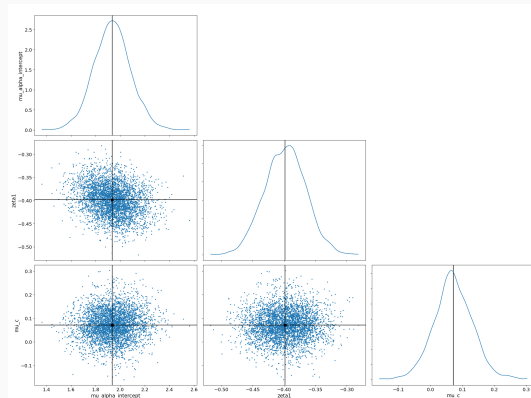
The identified model should exhibit good MCMC behavior.

The identified model should exhibit good MCMC behavior.

Expected Results:

  – Good convergence ($\hat{R} \approx 1$).

  – High Effective Sample Sizes (ESS).

  – Well-mixed, stable trace plots.

Now there should also be no problematic correlations involving 'mu_alpha_int'.

Key Points

Non-identifiability can arise in hierarchical models from redundant parameters (e.g., multiple intercepts).

Key Points

Non-identifiability can arise in hierarchical models from redundant parameters (e.g., multiple intercepts).

Additive structures like $d' = \text{person\_effect} + \text{condition\_effect}$ are common sources if both components have intercepts.

Key Points

Non-identifiability can arise in hierarchical models from redundant parameters (e.g., multiple intercepts).

Additive structures like $d' = \text{person\_effect} + \text{condition\_effect}$ are common sources if both components have intercepts.

Diagnose using MCMC convergence stats ($\hat{R}$, ESS) and visual checks (trace plots, pair plots).

Key Points

Non-identifiability can arise in hierarchical models from redundant parameters (e.g., multiple intercepts).

Additive structures like $d' = $ person_effect $+$ condition_effect are common sources if both components have intercepts.

Diagnose using MCMC convergence stats ($\hat{R}$, ESS) and visual checks (trace plots, pair plots).

Fix primarily via reparameterization (removing redundancy, using constraints).

Key Points

Non-identifiability can arise in hierarchical models from redundant parameters (e.g., multiple intercepts).

Additive structures like $d' = \text{person\_effect} + \text{condition\_effect}$ are common sources if both components have intercepts.

Diagnose using MCMC convergence stats ($\hat{R}$, ESS) and visual checks (trace plots, pair plots).

Fix primarily via reparameterization (removing redundancy, using constraints).

Always check for potential identifiability issues before interpreting model parameters. A model with convergence issues will generally not yield reliable inferences!

Identifiability

Joachim Vandekerckhove

Spring 2025