

Identifiability

Joachim Vandekerckhove

Spring 2025

Recap: What is Non-Identifiability?

A model is **non-identifiable** if different parameter combinations yield the exact same predictions (likelihood).

Recap: What is Non-Identifiability?

A model is **non-identifiable** if different parameter combinations yield the exact same predictions (likelihood).

The data cannot distinguish between these combinations.

Recap: What is Non-Identifiability?

A model is **non-identifiable** if different parameter combinations yield the exact same predictions (likelihood).

The data cannot distinguish between these combinations.

Structural Non-Identifiability: Inherent in model equations.

Recap: What is Non-Identifiability?

A model is **non-identifiable** if different parameter combinations yield the exact same predictions (likelihood).

The data cannot distinguish between these combinations.

Structural Non-Identifiability: Inherent in model equations.

Consequences:

- Poor MCMC convergence ($\hat{R} \gg 1$, low ESS).
- Unreliable posterior distributions (drifting, odd shapes).
- Strong correlations between confounded parameters.
- Misleading inferences.

Example: Hierarchical SDT Model

Predicting Hits (H_{ip}) and False Alarms (F_{ip}) for person p under condition i .

Example: Hierarchical SDT Model

Predicting Hits (H_{ip}) and False Alarms (F_{ip}) for person p under condition i .

Likelihood: Binomial counts based on SDT probabilities.

$$(H_{ip}, F_{ip}) \sim SDT(d'_{ip}, c_p, n_s, n_n)$$

Example: Hierarchical SDT Model

Predicting Hits (H_{ip}) and False Alarms (F_{ip}) for person p under condition i .

Likelihood: Binomial counts based on SDT probabilities.

$$(H_{ip}, F_{ip}) \sim SDT(d'_{ip}, c_p, n_s, n_n)$$

Hierarchical Parameter Models (Non-Identified Version):

$$d'_{ip} = \alpha_p + \gamma_i$$

$$\gamma_i = \zeta_0 + \zeta_1 \cdot \text{predictor}_i \quad (\text{Condition effect on } d')$$

$$\alpha_p \sim \text{Normal}(\mu_\alpha, \sigma_\alpha^2) \quad (\text{Person sensitivity intercept})$$

$$c_p \sim \text{Normal}(\mu_c, \sigma_c^2) \quad (\text{Person criterion})$$

Spotting the Non-Identifiability

The core issue lies in the definition of sensitivity d'_{ip} :

$$\begin{aligned}d'_{ip} &= \alpha_p + \gamma_i \\&= \alpha_p + (\zeta_0 + \zeta_1 \cdot \text{predictor}_i) \\&= (\alpha_p + \zeta_0) + \zeta_1 \cdot \text{predictor}_i\end{aligned}$$

Spotting the Non-Identifiability

The core issue lies in the definition of sensitivity d'_{ip} :

$$\begin{aligned}d'_{ip} &= \alpha_p + \gamma_i \\&= \alpha_p + (\zeta_0 + \zeta_1 \cdot \text{predictor}_i) \\&= (\alpha_p + \zeta_0) + \zeta_1 \cdot \text{predictor}_i\end{aligned}$$

The likelihood only depends on the sum $(\alpha_p + \zeta_0)$.

Spotting the Non-Identifiability

The core issue lies in the definition of sensitivity d'_{ip} :

$$\begin{aligned}d'_{ip} &= \alpha_p + \gamma_i \\&= \alpha_p + (\zeta_0 + \zeta_1 \cdot \text{predictor}_i) \\&= (\alpha_p + \zeta_0) + \zeta_1 \cdot \text{predictor}_i\end{aligned}$$

The likelihood only depends on the sum $(\alpha_p + \zeta_0)$.

Let $\tilde{\alpha}_p = \alpha_p - C$ and $\tilde{\zeta}_0 = \zeta_0 + C$.

Then $\tilde{\alpha}_p + \tilde{\zeta}_0 = \alpha_p + \zeta_0$. The predicted d'_{ip} is unchanged.

Spotting the Non-Identifiability

The core issue lies in the definition of sensitivity d'_{ip} :

$$\begin{aligned}d'_{ip} &= \alpha_p + \gamma_i \\&= \alpha_p + (\zeta_0 + \zeta_1 \cdot \text{predictor}_i) \\&= (\alpha_p + \zeta_0) + \zeta_1 \cdot \text{predictor}_i\end{aligned}$$

The likelihood only depends on the sum $(\alpha_p + \zeta_0)$.

Let $\tilde{\alpha}_p = \alpha_p - C$ and $\tilde{\zeta}_0 = \zeta_0 + C$.

Then $\tilde{\alpha}_p + \tilde{\zeta}_0 = \alpha_p + \zeta_0$. The predicted d'_{ip} is unchanged.

As a result, there cannot exist a unique set of optimal parameters.

Spotting the Non-Identifiability

The core issue lies in the definition of sensitivity d'_{ip} :

$$\begin{aligned}d'_{ip} &= \alpha_p + \gamma_i \\&= \alpha_p + (\zeta_0 + \zeta_1 \cdot \text{predictor}_i) \\&= (\alpha_p + \zeta_0) + \zeta_1 \cdot \text{predictor}_i\end{aligned}$$

The likelihood only depends on the sum $(\alpha_p + \zeta_0)$.

Let $\tilde{\alpha}_p = \alpha_p - C$ and $\tilde{\zeta}_0 = \zeta_0 + C$.

Then $\tilde{\alpha}_p + \tilde{\zeta}_0 = \alpha_p + \zeta_0$. The predicted d'_{ip} is unchanged.

As a result, there cannot exist a unique set of optimal parameters.

Nothing can distinguish trade-offs between μ_α and ζ_0 .

Spotting the Non-Identifiability

The core issue lies in the definition of sensitivity d'_{ip} :

$$\begin{aligned}d'_{ip} &= \alpha_p + \gamma_i \\&= \alpha_p + (\zeta_0 + \zeta_1 \cdot \text{predictor}_i) \\&= (\alpha_p + \zeta_0) + \zeta_1 \cdot \text{predictor}_i\end{aligned}$$

The likelihood only depends on the sum $(\alpha_p + \zeta_0)$.

Let $\tilde{\alpha}_p = \alpha_p - C$ and $\tilde{\zeta}_0 = \zeta_0 + C$.

Then $\tilde{\alpha}_p + \tilde{\zeta}_0 = \alpha_p + \zeta_0$. The predicted d'_{ip} is unchanged.

As a result, there cannot exist a unique set of optimal parameters.

Nothing can distinguish trade-offs between μ_α and ζ_0 .

This is **structural non-identifiability**.

Common MCMC Symptoms

Fitting the non-identified model is expected to show:

Common MCMC Symptoms

Fitting the non-identified model is expected to show:

- **Poor Convergence** for confounded parameters (μ_α, ζ_0) : High \hat{R} , Low ESS.

Common MCMC Symptoms

Fitting the non-identified model is expected to show:

- **Poor Convergence** for confounded parameters (μ_α, ζ_0): High \hat{R} , Low ESS.
- **Wandering Trace Plots**: Chains for μ_α, ζ_0 may not mix or stabilize.

Common MCMC Symptoms

Fitting the non-identified model is expected to show:

- **Poor Convergence** for confounded parameters (μ_α, ζ_0): High \hat{R} , Low ESS.
- **Wandering Trace Plots**: Chains for μ_α, ζ_0 may not mix or stabilize.
- **Strong Posterior Correlation**: A clear negative correlation between μ_α and ζ_0 samples.

Common MCMC Symptoms

Fitting the non-identified model is expected to show:

- **Poor Convergence** for confounded parameters (μ_α , ζ_0): High \hat{R} , Low ESS.
- **Wandering Trace Plots**: Chains for μ_α , ζ_0 may not mix or stabilize.
- **Strong Posterior Correlation**: A clear negative correlation between μ_α and ζ_0 samples.
- **Sampling Difficulties**: Potential warnings (divergences, etc.).

Any or all of these can occur.

PyMC Code (Non-Identified SDT)

Key parameter definitions causing the issue:

```
1 # Uninformative priors for Sensitivity (d') parameters
2 mu_alpha = pm.Normal("mu_alpha", mu=0.0, sigma=1.0e9) # << Part 1
3 sigma_alpha = pm.HalfNormal("sigma_alpha", sigma=1.0e9)
4 zeta0 = pm.Normal("zeta0", mu=0.0, sigma=1.0e9) # << Part 2
5 zeta1 = pm.Normal("zeta1", mu=0.0, sigma=1.0e9)
6
7 # Person sensitivity intercepts (drawn from mu_alpha)
8 alpha_p = pm.Normal("alpha_p", ...)
9
10 # Calculate d' using alpha_p and gamma_i
11 d_ip = pm.Deterministic("d_ip", alpha_p[...] + gamma_i) # << Part 3
12
13 # ... SDT calculations ...
```

Possible Diagnosis 1: Trace Plots

Use `'arviz.plot_trace()'` focusing on `'mu_alpha'` and `'zeta0'`.

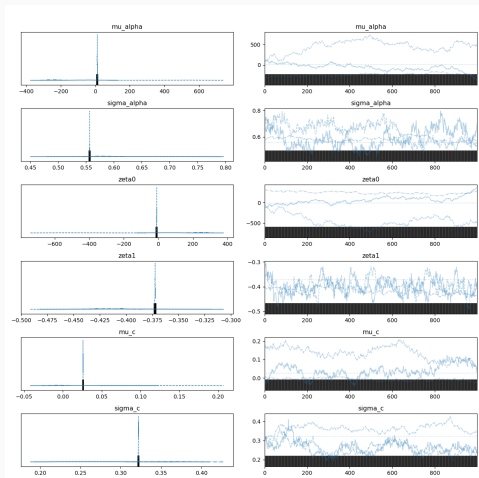
Possible Diagnosis 1: Trace Plots

Use `'arviz.plot_trace()'` focusing on `'mu_alpha'` and `'zeta0'`.

Look For:

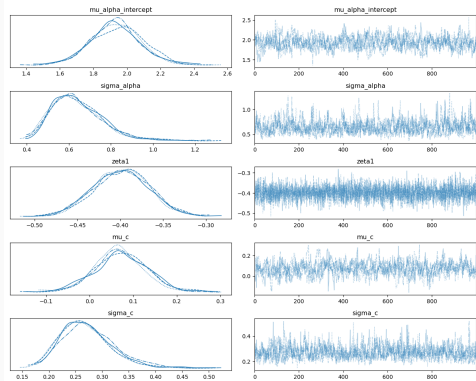
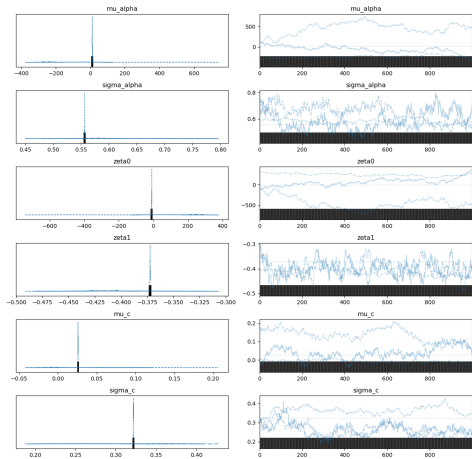
- Poor mixing for `'mu_alpha'`, `'zeta0'`.
- Chains exploring different levels.
- Compare to well-behaved traces (e.g., `'mu_c'`).
- Corroborate with high \hat{R} in summary.

Possible Diagnosis 1: Trace Plots



Possible Diagnosis 1: Trace Plots

For comparison, here are trace plots of the fixed (identified) model:



MCMC Output Table

--- Non-Identified SDT Model Summary ---

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_bulk	ess_tail	r_hat
mu_alpha	40.779	268.443	-304.985	570.827	128.271	70.199	5.0	17.0	2.70
zeta0	-38.809	268.503	-569.032	307.068	128.298	70.209	5.0	17.0	2.70
zeta1	-0.397	0.028	-0.446	-0.350	0.008	0.001	15.0	180.0	1.22
mu_c	0.048	0.058	-0.020	0.173	0.027	0.014	5.0	26.0	2.35

Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 379 seconds.
There were 1000 divergences after tuning. Increase 'target_accept' or reparameterize.

--- Identified SDT Model Summary ---

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_bulk	ess_tail	r_hat
mu_alpha_intercept	1.936	0.149	1.659	2.221	0.008	0.004	380.0	598.0	1.01
zeta1	-0.398	0.033	-0.463	-0.341	0.001	0.000	3558.0	3080.0	1.00
mu_c	0.072	0.062	-0.046	0.186	0.003	0.002	426.0	735.0	1.01

Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 12 seconds.

Possible Diagnosis 2: Pair Plots

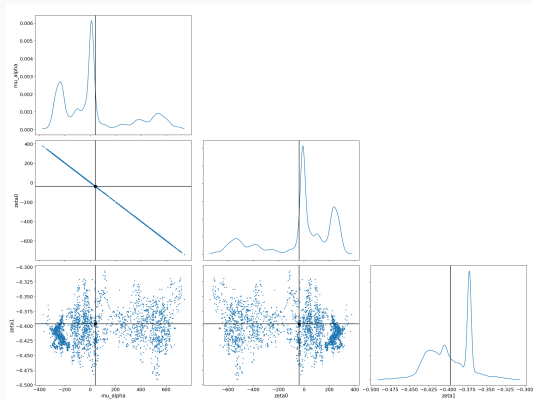
Use `'arviz.plot_pair()'` for sensitivity parameters `'mu_alpha'`, `'zeta0'`.

Possible Diagnosis 2: Pair Plots

Use `'arviz.plot_pair()'` for sensitivity parameters `'mu_alpha'`, `'zeta0'`.

Look For:

- Strong linear correlation “ridge.”
- Specifically, a **negative** correlation between `'mu_alpha'` and `'zeta0'`.
- Poorly conditioned marginal densities.



Fixing the Non-Identifiability

Common strategies:

Fixing the Non-Identifiability

Common strategies:

a) Reparameterization (Preferred):

- Remove redundant parameters (e.g., remove ζ_0).
- Use sum-to-zero constraints (more complex).
- Use contrasts relative to a baseline condition.

Fixing the Non-Identifiability

Common strategies:

a) Reparameterization (Preferred):

- Remove redundant parameters (e.g., remove ζ_0).
- Use sum-to-zero constraints (more complex).
- Use contrasts relative to a baseline condition.

b) Informative Priors: Can sometimes help practically, but remember that your conclusions will then strongly depend on your assumptions.

Fixing the Non-Identifiability

Common strategies:

a) Reparameterization (Preferred):

- Remove redundant parameters (e.g., remove ζ_0).
- Use sum-to-zero constraints (more complex).
- Use contrasts relative to a baseline condition.

b) Informative Priors: Can sometimes help practically, but remember that your conclusions will then strongly depend on your assumptions.

Let's see what informative priors can do...

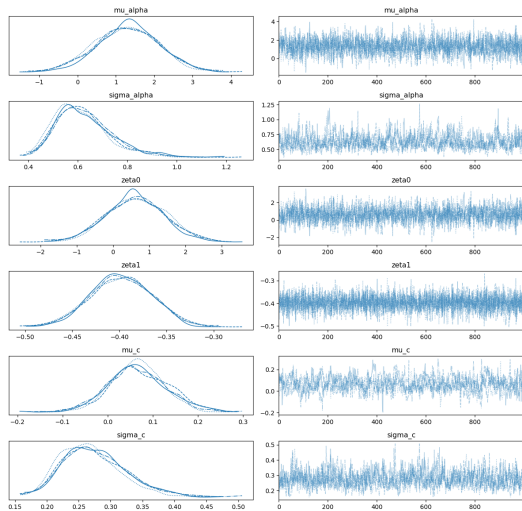
Using Informative Priors

Setting $\mu_\alpha \sim N(0, 1.5)$ gives converging chains:

```
--- Semi-Identified SDT Model Summary ---  
      mean      sd  hdi_3%  hdi_97%  mcse_mean  mcse_sd  ess_bulk  ess_tail  r_hat  
mu_alpha  1.301  0.827  -0.310   2.807    0.018    0.013   2023.0   2365.0   1.00  
zeta0     0.629  0.825  -0.924   2.183    0.018    0.013   2002.0   2339.0   1.00  
zeta1    -0.398  0.033  -0.459  -0.338    0.001    0.001   3915.0   2744.0   1.00  
mu_c      0.067  0.064  -0.051   0.190    0.003    0.002    480.0    720.0   1.01
```

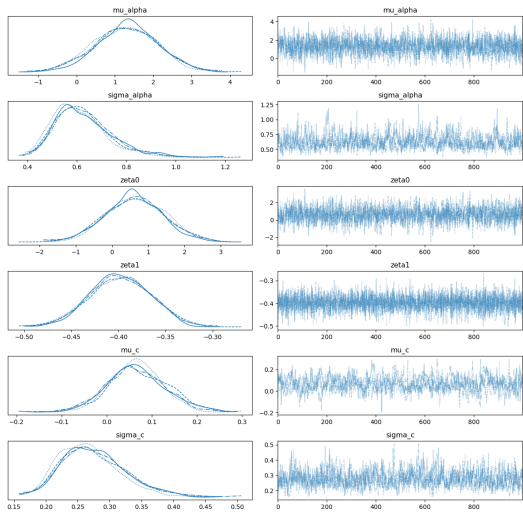

Using Informative Priors

Even reasonable trace plots:

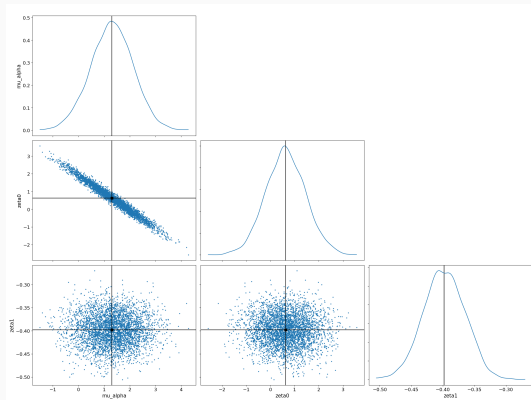


Using Informative Priors

Even reasonable trace plots:



But the trade-off pops out of the pair plot:



Applying the Real Fix: Remove ζ_0

Interpretation changes: μ_α becomes 'mu_alpha_int', average d' when predictor=0.

Applying the Real Fix: Remove ζ_0

Interpretation changes: μ_α becomes 'mu_alpha_int', average d' when predictor=0.

```
1 # Priors for Sensitivity (d') parameters
2 mu_alpha_int = pm.Normal("mu_alpha_int", ...) # Was mu_alpha
3 sigma_alpha = pm.HalfNormal("sigma_alpha", ...)
4 # zeta0 = pm.Normal("zeta0", ...) # <<< REMOVED
5 zeta1 = pm.Normal("zeta1", ...)
6
7 # Person sensitivity intercepts (drawn from mu_alpha_intercept)
8 alpha_p = pm.Normal("alpha_p", ...)
9
10 # Condition Effects on Sensitivity (NO intercept)
11 gamma_i = pm.Deterministic("gamma_i", zeta1 * cond_pred_data)
12
13 # SDT likelihood calculations are structurally identical
14 # ...
```

Results After Fixing

The identified model should exhibit good MCMC behavior.

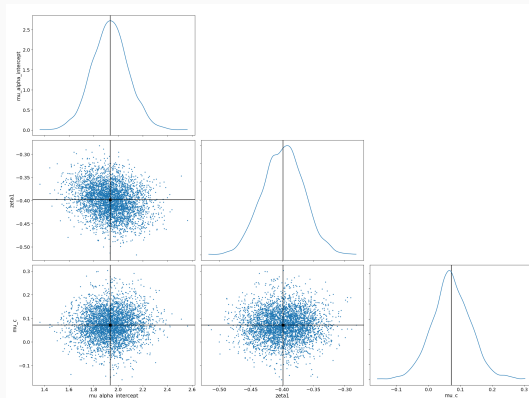
Results After Fixing

The identified model should exhibit good MCMC behavior.

Expected Results:

- Good convergence ($\hat{R} \approx 1$).
- High Effective Sample Sizes (ESS).
- Well-mixed, stable trace plots.

Now there should also be no problematic correlations involving 'mu_alpha_int'.



Key Points

Non-identifiability can arise in hierarchical models from redundant parameters (e.g., multiple intercepts).

Key Points

Non-identifiability can arise in hierarchical models from redundant parameters (e.g., multiple intercepts).

Additive structures like $d' = \text{person_effect} + \text{condition_effect}$ are common sources if both components have intercepts.

Key Points

Non-identifiability can arise in hierarchical models from redundant parameters (e.g., multiple intercepts).

Additive structures like $d' = \text{person_effect} + \text{condition_effect}$ are common sources if both components have intercepts.

Diagnose using MCMC convergence stats (\hat{R} , ESS) and visual checks (trace plots, pair plots).

Key Points

Non-identifiability can arise in hierarchical models from redundant parameters (e.g., multiple intercepts).

Additive structures like $d' = \text{person_effect} + \text{condition_effect}$ are common sources if both components have intercepts.

Diagnose using MCMC convergence stats (\hat{R} , ESS) and visual checks (trace plots, pair plots).

Fix primarily via reparameterization (removing redundancy, using constraints).

Key Points

Non-identifiability can arise in hierarchical models from redundant parameters (e.g., multiple intercepts).

Additive structures like $d' = \text{person_effect} + \text{condition_effect}$ are common sources if both components have intercepts.

Diagnose using MCMC convergence stats (\hat{R} , ESS) and visual checks (trace plots, pair plots).

Fix primarily via reparameterization (removing redundancy, using constraints).

Always check for potential identifiability issues before interpreting model parameters. A model with convergence issues will generally not yield reliable inferences!

Identifiability

Joachim Vandekerckhove

Spring 2025