

# Identifiability

In statistical and cognitive modeling, particularly when we deal with complex or custom models, a crucial concept is *identifiability*. A model is considered *non-identifiable* if there exist different combinations of the model's parameters that always produce exactly the same predictions. In the context of likelihood-based models, this implies that distinct parameter sets yield identical likelihood values for any potential observed data. As a result no amount of observed data can distinguish between these parameter constellations. This makes it impossible to infer parameters from data – the data do not provide enough information to uniquely determine the true underlying parameters.

There are nuances to different kinds of non-identifiability, but certainly the most problematic type is **structural non-identifiability**, which is not due to insufficient data or poor experimental design, but is inherent in the mathematical structure of the model's equations themselves.

Fitting a non-identifiable model using MCMC methods typically results in several observable difficulties:

- Poor MCMC convergence, characterized by high values of the  $\hat{R}$  convergence statistic (much greater than 1) and low Effective Sample Sizes (ESS).
- Posterior distributions that look unstable, that drift, or take on unusual shapes.
- Strong correlations among the parameters affected by the non-identifiability in the posterior samples.

Those three symptoms are visible – but the fourth symptom is the problem: unidentified models are not suitable for statistical inference. You cannot estimate their parameters and you cannot draw conclusions about them. Avoiding or addressing non-identifiability is a necessary step for valid model-based inference.

## Illustrative Example: A Hierarchical Signal Detection Theory Model

Consider a hierarchical Signal Detection Theory (SDT) model as an example. Consider a simple SDT model that predicts counts of Hits ( $H_{ip}$ ) and False Alarms ( $F_{ip}$ ) for person  $p$  in condition  $i$ .

The data are modeled via a **likelihood** derived from Binomial distributions, where the success probabilities are determined by SDT parameters: sensitivity ( $d'_{ip}$ ) and criterion ( $c_p$ ), along with trial counts ( $n_s, n_n$ ). Jointly, these equations (which we have covered many times) can be used to compute the likelihood of the data given the parameters. We can write this abstractly as:

$$(H_{ip}, F_{ip}) \sim SDT(d'_{ip}, c_p, n_s, n_n)$$

Within a **hierarchical signal detection theory (HSDT) model**, the parameters  $d'_{ip}$  and  $c_p$  have some sort of structure. Let's suppose a very reasonable HSDT model, in which  $d'_{ip}$  is defined as the sum of a person-specific intercept  $\alpha_p$  and a condition-specific effect  $\gamma_i$ :

$$d'_{ip} = \alpha_p + \gamma_i$$

The condition effect  $\gamma_i$  is specified linearly based on a condition-level predictor:

$$\gamma_i = \zeta_0 + \zeta_1 \cdot \text{predictor}_i \quad (\text{Condition effect on } d')$$

$\zeta_0$  functions as the intercept for the condition effect and  $\zeta_1$  is the slope (i.e., effect size) of the condition effect.

The person intercept  $\alpha_p$  is drawn from a group-level distribution to allow pooling of information across people:

$$\alpha_p \sim \text{Normal}(\mu_\alpha, \sigma_\alpha^2) \quad (\text{Person sensitivity intercept})$$

The person criterion  $c_p$  is also modeled hierarchically (but less relevant here):

$$c_p \sim \text{Normal}(\mu_c, \sigma_c^2) \quad (\text{Person criterion})$$

The identifiability issue in this formulation is concentrated within the specification of  $d'_{ip}$ .

## Analysis of Non-Identifiability

The source of the non-identifiability in this model is revealed by substituting the definition of  $\gamma_i$  into the equation for  $d'_{ip}$ :

$$\begin{aligned} d'_{ip} &= \alpha_p + \gamma_i \\ &= \alpha_p + (\zeta_0 + \zeta_1 \cdot \text{predictor}_i) \\ &= (\alpha_p + \zeta_0) + \zeta_1 \cdot \text{predictor}_i \end{aligned}$$

The parameter  $d'_{ip}$  depends on the sum of the person intercept  $\alpha_p$  and the condition effect intercept  $\zeta_0$ , in addition to the term involving  $\zeta_1$  and the predictor. A jargon-y way of describing this setup is that  $d'_{ip}$  has two intercepts (or rather, two parameters that each act as an intercept).

Since the model likelihood depends on  $d'_{ip}$  and  $c_p$ , it is determined by the value of the sum  $(\alpha_p + \zeta_0)$  and the term  $\zeta_1 \cdot \text{predictor}_i$ .

We can now consider two distinct parameter sets. Let one set contain  $(\alpha_p, \zeta_0)$  and another contain  $(\tilde{\alpha}_p, \tilde{\zeta}_0)$ , where  $\tilde{\alpha}_p = \alpha_p - C$  and  $\tilde{\zeta}_0 = \zeta_0 + C$  for an arbitrary constant  $C$ . The sum for the second set is:

$$\tilde{\alpha}_p + \tilde{\zeta}_0 = (\alpha_p - C) + (\zeta_0 + C) = \alpha_p + \zeta_0$$

As is hopefully obvious, any constant  $C$  could be added to  $\zeta_0$  and subtracted from  $\alpha_p$  without altering their sum. Consequently, the calculated value of  $d'_{ip}$  remains unchanged. As  $d'_{ip}$  (and an identifiable  $c_p$ ) determines the likelihood, these distinct parameter combinations yield identical likelihood distributions, always. In other words, the individual values of  $\alpha_p$  and  $\zeta_0$  are arbitrary (even though their sum is not).

$\alpha_p$  and  $\zeta_0$  are said to be confounded additively. At the group level, this manifests as an inability to distinguish between the mean person sensitivity  $\mu_\alpha$  and the condition intercept  $\zeta_0$ . The model structure permits trade-offs between these parameters while maintaining the same likelihood. This constitutes *structural non-identifiability*.

## Common MCMC Symptoms

MCMC diagnostics can help us identify non-identifiability. When a non-identified model is fit using MCMC methods, several diagnostic indicators should raise red flags:

- **Poor Convergence** for the parameters involved in the confounding (e.g.,  $\mu_\alpha$  and  $\zeta_0$ ). This is shown by high  $\hat{R}$  values and low Effective Sample Sizes (ESS).
- **Wandering Trace Plots.** The MCMC chains for the confounded parameters may exhibit poor mixing and fail to stabilize, often appearing to explore a ‘ridge’ of high likelihood.
- **Strong Posterior Correlation.** Visual analysis of the joint posterior distribution (or pairwise plots) of the confounded parameters may show a dependency. In this additive case, a negative linear correlation between  $\mu_\alpha$  and  $\zeta_0$  is expected, but in general the relationship need not be linear, and it may even involve many parameters. (In fact, in the example scenario, every  $\alpha_p$  is also involved.)
- **Sampling Difficulties.** Numerical stability issues may arise, potentially leading to sampler warnings such as divergences.

These observations collectively signal the presence of non-identifiability.

## PyMC Implementation of the Non-Identified Model

The following PyMC code snippet illustrates the parameter definitions that lead to the non-identifiability in the hierarchical SDT model. (Find the full code in the GitHub repository, 0-introduction/src/sdt/sdt\_identifiability.py.)

```
1 # Uninformative priors for Sensitivity (d') parameters
2 mu_alpha = pm.Normal("mu_alpha", mu=0.0, sigma=1.0e9) # << Part 1
3 sigma_alpha = pm.HalfNormal("sigma_alpha", sigma=1.0e9)
4 zeta0 = pm.Normal("zeta0", mu=0.0, sigma=1.0e9) # << Part 2
5 zeta1 = pm.Normal("zeta1", mu=0.0, sigma=1.0e9)
6
7 # Person sensitivity intercepts (drawn from mu_alpha)
8 alpha_p = pm.Normal("alpha_p", mu=mu_alpha, sigma=sigma_alpha)
9
10 # Calculate gamma_i (Condition effect on d')
11 # (Assuming 'cond_pred_data' is available elsewhere)
12 gamma_i = zeta0 + zeta1 * cond_pred_data
13
14 # Calculate d' using alpha_p and gamma_i
15 d_ip = pm.Deterministic("d_ip", alpha_p[...] + gamma_i) # << Part 3
16
17 # ... SDT likelihood calculations using d_ip and c_p ...
```

Listing 1: PyMC code snippet for the non-identified SDT model parameters.

The inclusion of ‘mu\_alpha’ (Part 1, mean of person intercepts) and ‘zeta0’ (Part 2, intercept of condition effects), combined with their additive contribution to ‘d\_ip’ (Part 3 via ‘alpha\_p’ and ‘gamma\_i’), introduces a structural redundancy.

## Diagnosis Method 1: Trace Plots

Evaluation of MCMC trace plots provides a visual diagnostic for convergence and mixing issues. Using `arviz.plot_trace()`, focused on parameters such as `mu_alpha` and `zeta0`, is standard practice.

When reviewing trace plots, pay attention to:

- Indications of poor mixing among the distinct chains for `mu_alpha` and `zeta0` – the chains should look like they are drawing samples from a stationary distribution and not wandering up and down. A well-mixed trace plot is sometimes described as a ‘fat, hairy caterpillar’.
- Evidence of chains exploring separate regions or failing to converge to a common stationary distribution. All chains should converge to the same region of the parameter space.
- Cross-check visual findings with high  $\hat{R}$  values reported in the MCMC summary.

Figure 1 presents trace plots from a fit of the non-identified model. Here we see the expected poor mixing and wandering for the confounded parameters. By contrast, Figure 2 shows well-behaved traces from an identified version of the model.

## MCMC Output Summary

Quantitative assessment of MCMC performance is provided by the summary table. The summary for the non-identified model looks like this:

```
--- Non-Identified SDT Model Summary ---
      mean      sd   hdi_3%   hdi_97%  mcse_mean  mcse_sd  ess_bulk  ess_tail  r_hat
mu_alpha  40.779 268.443 -304.985  570.827   128.271   70.199     5.0    17.0    2.70
zeta0     -38.809 268.503 -569.032  307.068   128.298   70.209     5.0    17.0    2.70
zeta1      -0.397  0.028  -0.446  -0.350     0.008    0.001    15.0   180.0    1.22
mu_c        0.048  0.058  -0.020   0.173     0.027    0.014     5.0    26.0    2.35
Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 379s.
There were 1000 divergences after tuning. Increase `target_accept` or reparameterize.
```

The high  $\hat{R}$  values (2.70) and extremely low ESS (5.0) for `mu_alpha` and `zeta0` indicate poor convergence. Seeing 1000 divergences is also a red flag (ideally, there are none).

For comparison, the summary from the identified model below indicates successful convergence, with  $\hat{R}$  values near 1.0 and substantially higher ESS for all parameters, including the newly named `mu_alpha_intercept`.

```
--- Identified SDT Model Summary ---
      mean      sd   hdi_3%   hdi_97%  mcse_mean  mcse_sd  ess_bulk  ess_tail  r_hat
mu_alpha_int  1.936  0.149   1.659   2.221     0.008    0.004   380.0   598.0    1.01
zeta1        -0.398  0.033  -0.463  -0.341     0.001    0.000  3558.0  3080.0    1.00
mu_c          0.072  0.062  -0.046   0.186     0.003    0.002   426.0   735.0    1.01
Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 12s.
```

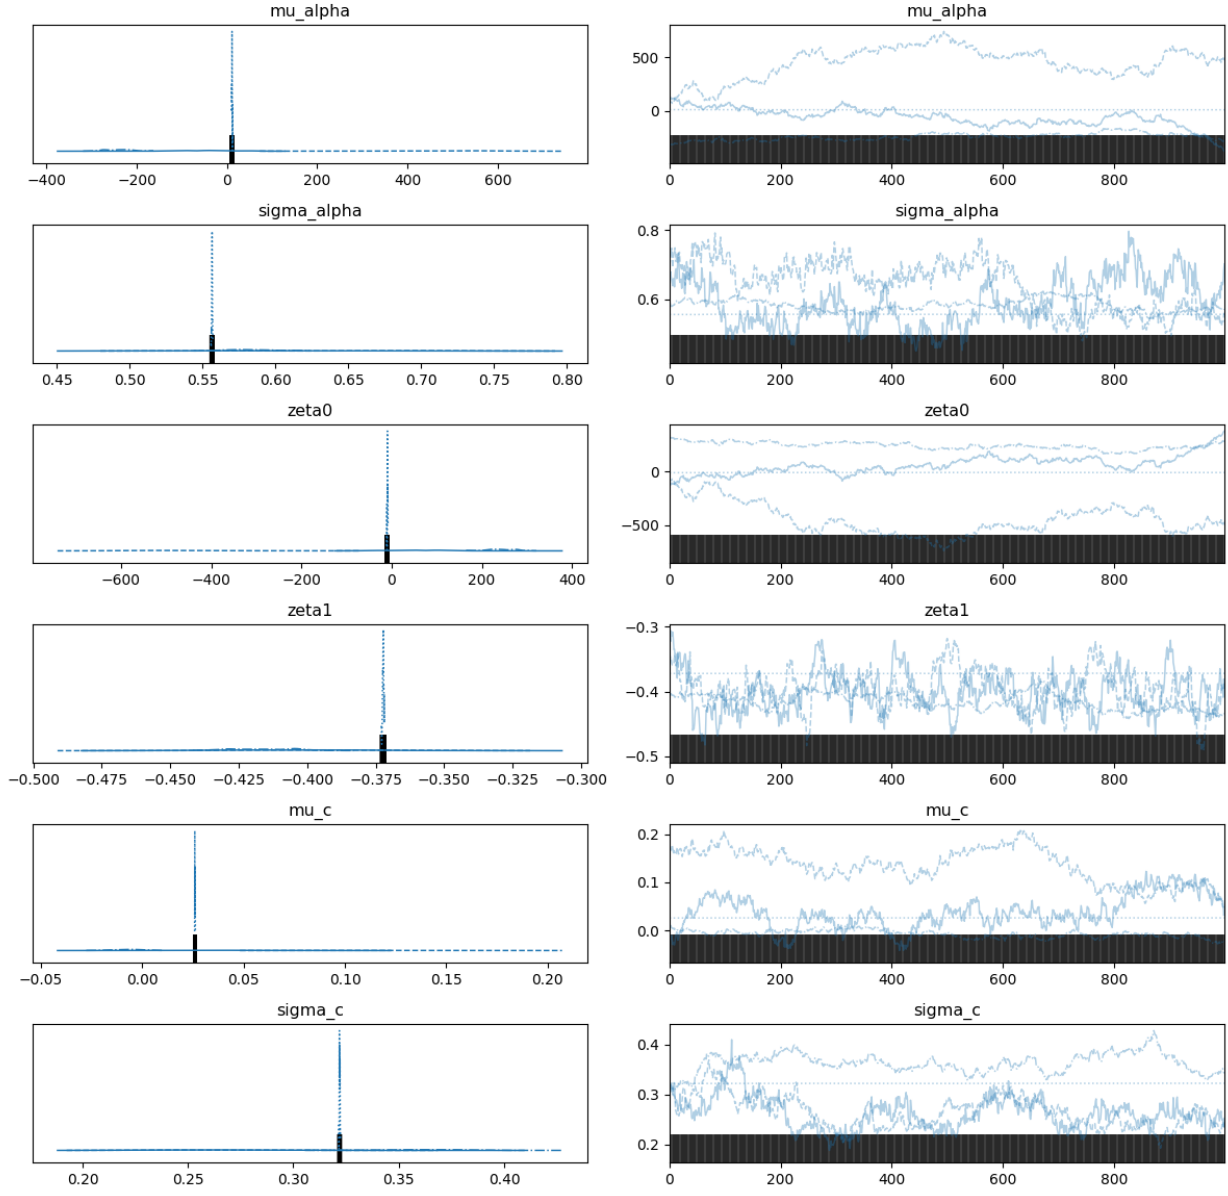


Figure 1: Trace plots for key parameters from the non-identified SDT model. Poor mixing and wandering chains are  $\mu_\alpha$  and  $\zeta_0$ 's lot.

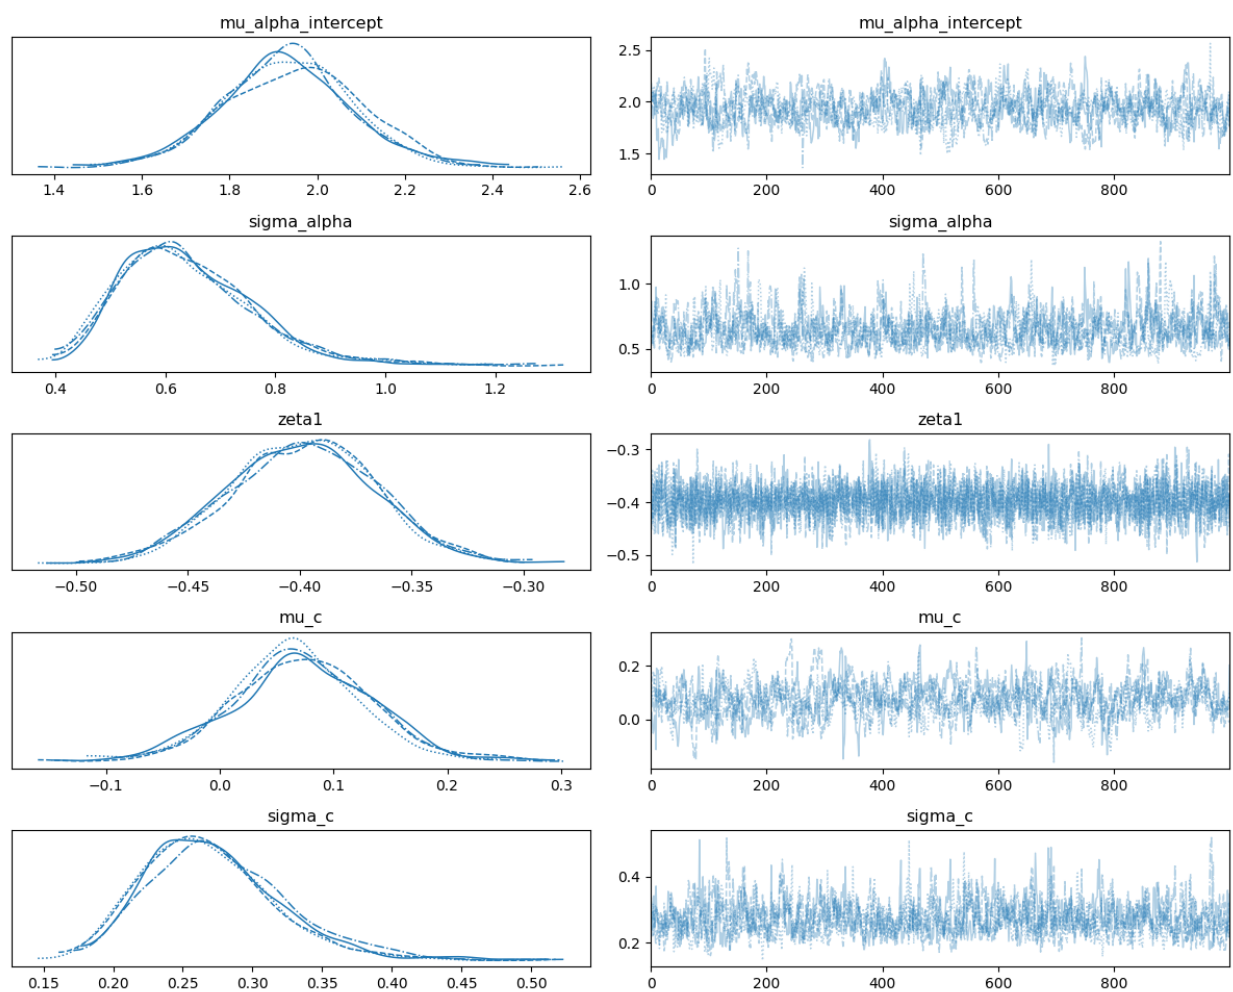


Figure 2: Trace plots for key parameters from the identified SDT model. Chains show good mixing and convergence.

## Diagnosis Method 2: Pair Plots

Pair plots show the joint posterior distributions of pairs of parameters. This is a diagnostic for (possibly non-linear) correlations indicative of identifiability issues. The `arviz.plot_pair()` function, applied to `mu_alpha` and `zeta0`, is useful here.

When examining pair plots for potentially confounded parameters, features to identify include:

- A distinct correlation forming a “ridge” across the scatter plot of posterior samples.
- Specifically for this additive confounding, a *negative, linear* correlation between `mu_alpha` and `zeta0`.
- Marginal distributions along the diagonal that may appear unusually flat or poorly defined – they will have excess variance compared the true posterior distributions.

Figure 3 displays the pair plot for `mu_alpha` and `zeta0` from the non-identified model – a strong negative linear correlation is visible.

## Strategies for Addressing Non-Identifiability

Non-identifiability must be addressed to maintain valid statistical inference. Common strategies include:

- a) Reparameterization (Preferred).** This involves modifying the model’s parameterization to remove the inherent redundancy. Methods include:
  - Removing redundant parameters. In the example, removing  $\zeta_0$  eliminates the additive confounding.
  - Imposing sum-to-zero constraints on effects. This is an alternative method to resolve additive confounding but a little trickier to implement.
  - Modeling effects as contrasts relative to a baseline condition. This implicitly removes a redundant intercept.

Reparameterization addresses the structural source of the issue.

- b) Informative Priors.** Applying informative priors to confounded parameters can constrain the posterior and improve MCMC convergence diagnostics. However, this does not resolve the *structural* non-identifiability. The resulting posterior distributions for the affected parameters will be heavily influenced by the chosen prior. That is not by itself a problem, but it is important not to forget that you may have made prior assumptions that drive your conclusions.
- c) Exogenous Predictors.** Adding an exogenous predictor that is not confounded with the parameters of interest can resolve the non-identifiability. In the example, replacing the mean person intercept  $\mu_\alpha$  with a function of a predictor variable (e.g.,  $\mu_\alpha = \zeta_1 \cdot \text{predictor}_i$ ) removes the confounding.

Even though this is not the preferred approach, let’s briefly look at the effect of informative priors.

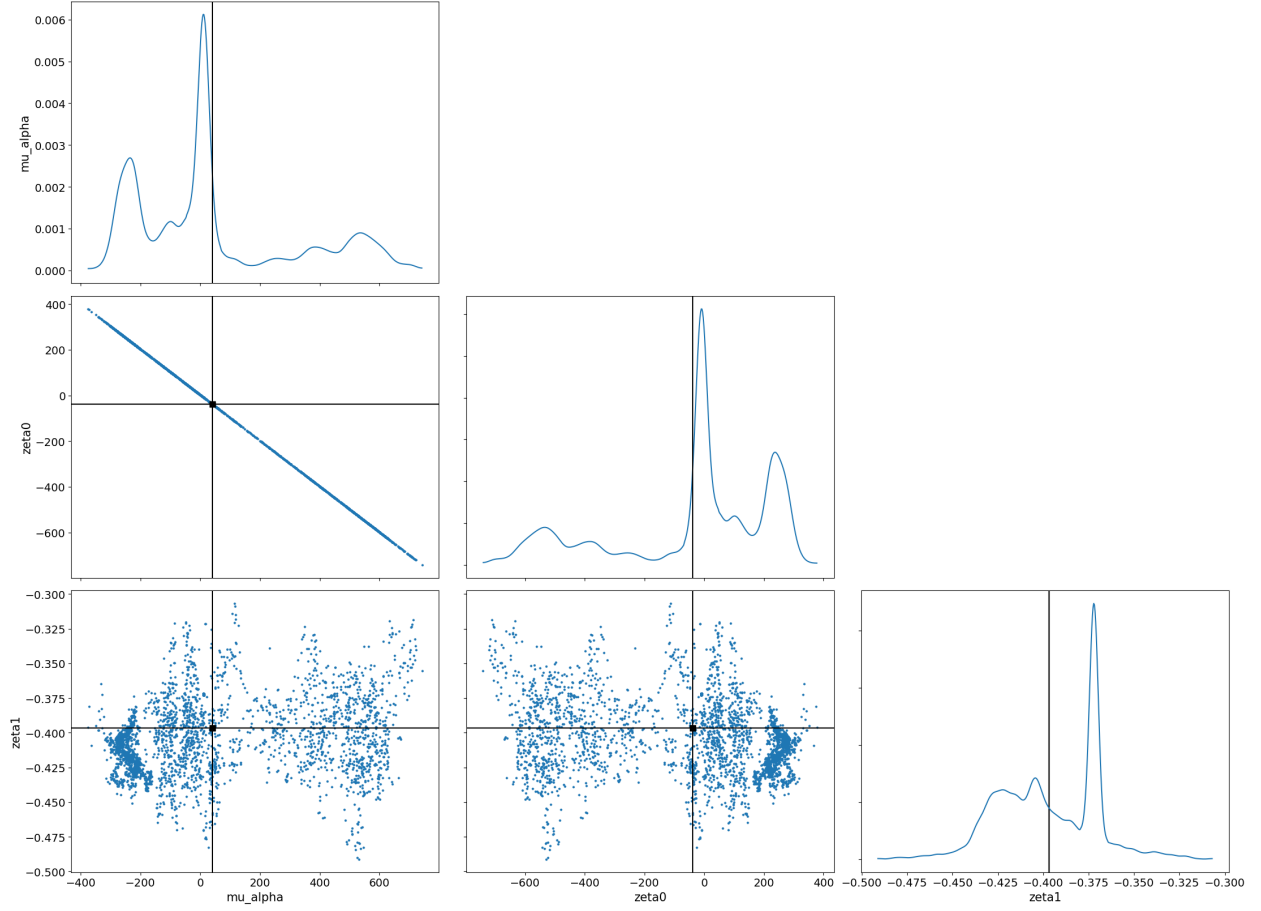


Figure 3: Pair plot for  $\mu_{\alpha}$  and  $\zeta_0$  from the non-identified SDT model. A strong negative linear correlation indicates confounding, and the marginal posteriors on the diagonal have strange complex shapes.



## Effect of Informative Priors

Employing informative priors can lead to improved MCMC convergence statistics even in a structurally non-identified model. I sometimes call those models “classically unidentified” because they are not identified by the data, but they can be identified by the data and a Bayesian prior.

Using  $\mu_\alpha \sim \text{Normal}(0, 1.5)$  (a more informative prior than before) in the non-identified SDT model yields convergence statistics shown below.

--- Semi-Identified SDT Model Summary ---									
mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_bulk	ess_tail	r_hat	
mu_alpha	1.301	0.827	-0.310	2.807	0.018	0.013	2023.0	2365.0	1.00
zeta0	0.629	0.825	-0.924	2.183	0.018	0.013	2002.0	2339.0	1.00
zeta1	-0.398	0.033	-0.459	-0.338	0.001	0.001	3915.0	2744.0	1.00
mu_c	0.067	0.064	-0.051	0.190	0.003	0.002	480.0	720.0	1.01

The  $\hat{R}$  values are now near 1.00, and ESS values are considerably higher than in the standard non-identified fit. Trace plots (Figure 4) also appear more stable. This looks almost good!

However, the pair plot for ‘mu\_alpha’ and ‘zeta0’ (Figure 5) continues to show the underlying negative correlation ridge, even if it is a little jittered now. While the prior aided convergence, the structural confounding remains.

## Implementation of the Structural Fix: Reparameterization

The best approach to resolve the structural non-identifiability in this SDT model is reparameterization by removing the condition effect intercept  $\zeta_0$ .

This adjustment necessitates a change in parameter interpretation. The parameter previously known as  $\mu_\alpha$  now represents the average  $d'$  when the predictor variable is zero. This parameter is renamed ‘mu\_alpha\_intercept’ to reflect this. The parameter  $\zeta_1$  continues to represent the rate of change in  $d'$  with respect to the predictor, now relative to this new baseline.

The modified PyMC code for the identified model is as follows:

```

1 # Priors for Sensitivity (d') parameters
2 mu_alpha_int = pm.Normal("mu_alpha_int", mu=0.0, sigma=1.0e9) # Was mu_alpha
3 sigma_alpha = pm.HalfNormal("sigma_alpha", sigma=1.0e9)
4 # zeta0 = pm.Normal("zeta0", mu=0.0, sigma=1.0e9) # <<< REMOVED
5 zeta1 = pm.Normal("zeta1", mu=0.0, sigma=1.0e9)
6
7 # Person sensitivity intercepts (drawn from mu_alpha_intercept)
8 alpha_p = pm.Normal("alpha_p", mu=mu_alpha_int, sigma=sigma_alpha)
9
10 # Condition Effects on Sensitivity (NO intercept)
11 gamma_i = pm.Deterministic("gamma_i", zeta1 * cond_pred_data)
12
13 # Calculate d' using alpha_p and gamma_i
14 d_ip = pm.Deterministic("d_ip", alpha_p[...] + gamma_i)
15
16 # SDT likelihood calculations are structurally identical

```

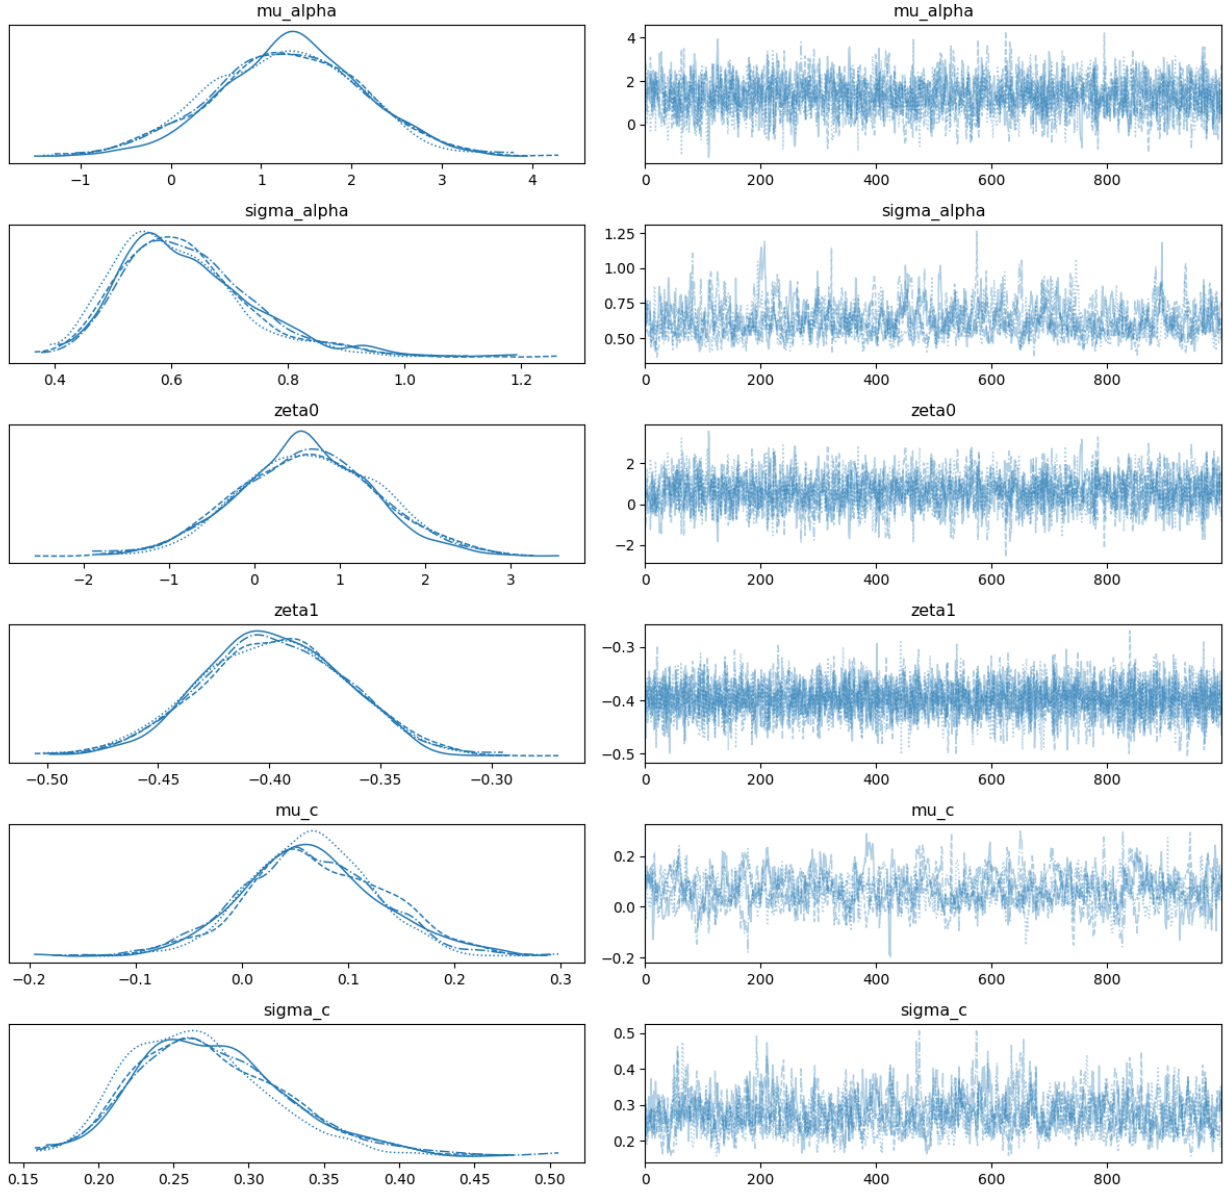


Figure 4: Trace plots from the model with an informative prior on  $\mu_\alpha$ . Apparent convergence based on traces.

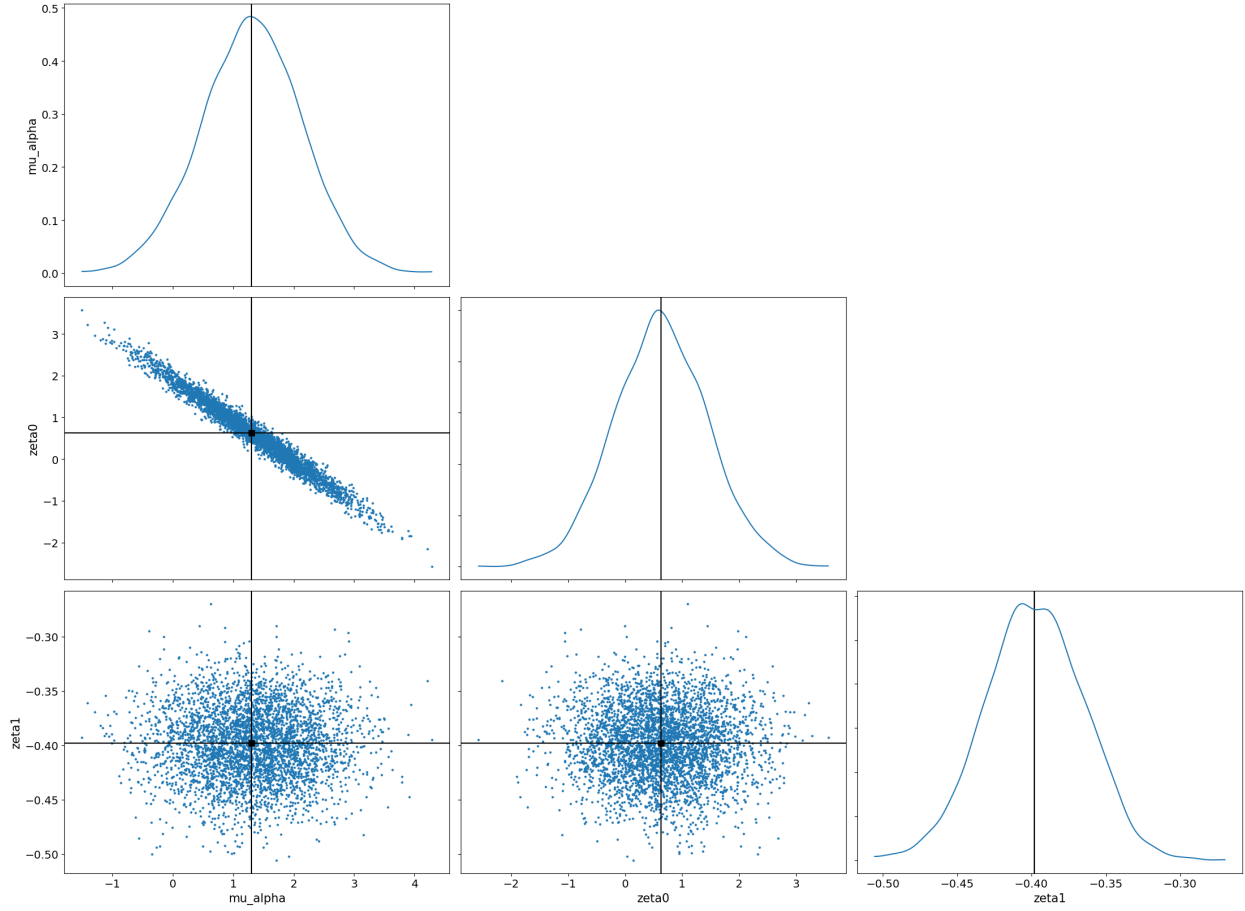


Figure 5: Pair plot for  $\mu_\alpha$  and  $\zeta_0$  from the model with informative prior. The negative correlation ridge persists.

17 # ...

Listing 2: PyMC code snippet for the identified SDT model parameters (removing  $\zeta_0$ ).

By removing the ‘zeta0’ term and defining ‘gamma\_i’ solely based on ‘zeta1’ and the predictor, the additive confounding is eliminated, resulting in a structurally identified model.

## Results Following Reparameterization

If we now fit the reparameterized, identified model using MCMC, we see that the convergence and diagnostics are improved. (Table repeated from before.)

```

--- Identified SDT Model Summary ---
              mean      sd   hdi_3%   hdi_97%   mcse_mean   mcse_sd   ess_bulk   ess_tail   r_hat
mu_alpha_int  1.936    0.149    1.659    2.221      0.008      0.004     380.0     598.0     1.01
zeta1         -0.398    0.033   -0.463   -0.341      0.001      0.000    3558.0    3080.0     1.00
mu_c           0.072    0.062   -0.046    0.186      0.003      0.002     426.0     735.0     1.01
Sampling 4 chains for 1_500 tune and 1_000 draw iterations (6_000 + 4_000 draws total) took 12s.

```

We see:

- Good convergence ( $\hat{R} \approx 1$ ) for all parameters.
- High Effective Sample Sizes (ESS).
- Well-mixed, stable trace plots (in Figure 2).

Following the removal of the confounding source, problematic correlations involving the parameter ‘mu\_alpha\_intercept’ or other model parameters are gone.

Figure 6 shows the pair plot for key parameters from the identified model, including ‘mu\_alpha\_intercept’. The absence of strong linear correlations shows that the structural non-identifiability has been successfully resolved.

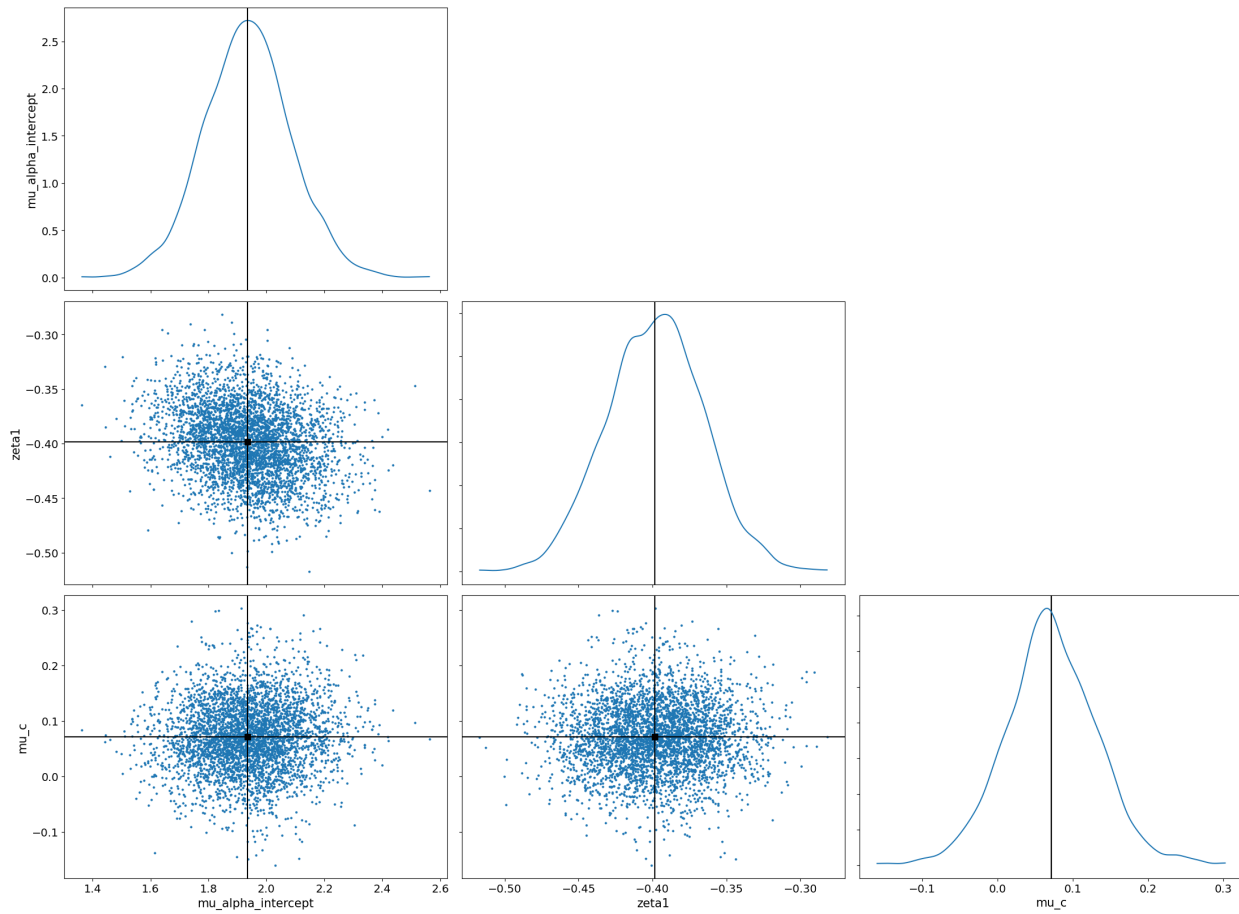


Figure 6: Pair plot for key parameters from the identified SDT model. No strong correlations are evident.

## Summary Points

- Non-identifiability is a condition where distinct parameter sets produce identical data likelihoods, preventing unique parameter estimation.

- In hierarchical models, non-identifiability can arise from redundant parameters, commonly seen with additive intercepts in model components.
- The structure  $d' = \text{person\_effect} + \text{condition\_effect}$  leads to non-identifiability if both effects incorporate their own intercept terms.
- Diagnosis relies on MCMC convergence statistics ( $\hat{R}$ , ESS) and visual checks such as trace plots and pair plots.
- The recommended solution for structural non-identifiability is reparameterization, which removes the underlying redundancy. While informative priors can sometimes assist convergence, they do not fundamentally resolve the structural issue.
- It is imperative to assess potential identifiability issues before interpreting model parameters. Models exhibiting convergence problems due to non-identifiability are unlikely to provide reliable inferences regarding the specific parameter values.